Microsoft 365 Copilot Analytics

Multi-Agency Architecture Decision Guide

Choose the right architecture when multiple agencies share one M365 tenant

Document Property	Value
Layer	1 — Strategy
Audience	Solution Architect, IT Admin, CIO
Version	4.0
Published	April 2026

> **How This Document Fits**
>
> This is a **Layer 1 — Strategy** document. Read this to choose between the two multi-agency architectures. Then hand off to IT with the Multi-Agency Setup Guide for step-by-step implementation.

1. When You Need This Guide

This guide applies when your organisation meets **all** of the following:

- **Single M365 tenant** — all agencies share one tenant
- **Multiple independent agencies** (or business units, divisions, departments) that operate semi-autonomously
- **Each agency needs Copilot analytics** scoped to their own population

Two architectural options exist. The right choice depends on your data visibility requirements and how much implementation effort you can invest.

Single-Agency Tenants

If your tenant serves only one agency (or your agencies don’t need separate analytics scoping), skip this guide entirely. Follow the standard Admin Setup Guide instead.

2. Decision Framework

The table below provides the key differences at a glance. Read both columns, then use the decision tree in Section 3 to confirm your choice.

Dimension	Default Priority Architecture
Core principle	Each agency CEO sees their team's data by default via the *"Your group"* scope on the Co
Cross-agency visibility	Acceptable — users with sufficient hierarchy access *can* navigate to "Your company" or o
How scoping works	Native AutoCxO identification + delegation + Entra ID manager hierarchy
Partitions needed	No — single default partition is sufficient
Report delivery	Native Copilot Dashboard ("Your group" scope)
PowerShell required	No
Typical setup effort	2–4 hours
Best for	Organisations where cross-agency data visibility is acceptable and speed of deployment i

3. Decision Tree

Walk through these questions in order. Each answer narrows your architecture and data path.

Question 1 — Must agencies be prevented from seeing each other's data?

Answer	Result
No — cross-agency visibility is acceptable	→ **Default Priority Architecture**. Proceed to Section 4, Default Priority
Yes — hard data isolation is required	→ **Isolation Architecture**. Continue to Question 2.

Question 2 — Do agencies share one email domain?

Only relevant if you chose Isolation Architecture above.

Answer	Result
No — each agency has its own email domain (e.g., `@agencyA.gov`, `@agencyB.gov`)	→ **Isolation Path B** (zero CSV
Yes — all agencies share one domain (e.g., `@government.gov`)	→ **Isolation Path A** (CSV requ

Question 3 — Do you need custom attribute filtering?

Applies to both Path A and Path B.

Answer	Result
Yes — you want to filter by Job Function, Location, Cost Centre, or other attributes	→ CSV org data upload is requi
No — partitioning by agency is sufficient	→ Path A: CSV with `Agency` at

Summary of Paths

Path	Shared Domain?	CSV Required?	Partition Filter Attribute
Default Priority	N/A	Optional (enrichment only)	None — uses manager hierarchy
Isolation Path A	Yes	**Yes** (with `Agency` column)	Custom `Agency` attribute
Isolation Path B	No	Only if custom attributes needed	System-generated `Domain` attribute

4. Data Strategy for Each Architecture

4.1 Default Priority Architecture

Org data is **optional enrichment** in this architecture. The Copilot Dashboard scoping is handled entirely by the Entra ID manager hierarchy — AutoCxO identifies the top leaders, and each agency CEO sees "Your group" by default.

- **Without org data:** Dashboard works with Department filter only (from Entra ID). No Location, Job Function, or custom attribute filters.
- **With org data:** All additional filters become available. Useful if agencies want richer analytics, but not required for the architecture to function.

4.2 Isolation Path A — CSV with Agency Attribute

When all agencies share one email domain, there is no system-level attribute to distinguish them. You must:

1. **Prepare a CSV file** with at least `Microsoft_PersonEmail`, `Microsoft_ManagerEmail`, and a custom `Agency` attribute
2. **Upload monthly** via the Viva Insights admin portal (or automate via the advanced data import API)
3. **Create one partition per agency** using the `Agency` attribute as the filter
4. **Assign an analyst** to each partition to run queries and build Power BI reports

> **Monthly Re-Upload**
>
> Org data is not automatically refreshed. Establish a monthly cadence to re-upload the CSV with current employee data. Stale data leads to inaccurate partition population and analytics.

4.3 Isolation Path B — Domain Attribute (Zero CSV)

When each agency has its own email domain, Viva Insights automatically generates a `Domain` attribute from users' email addresses. This requires **no CSV upload** for partitioning:

1. **Verify** that each agency's users have distinct email domains in Entra ID
2. **Create one partition per agency** using the system-generated `Domain` attribute as the filter (e.g., `Domain = agencyA.gov`)
3. **Assign an analyst** to each partition

If you later need custom attributes (Job Function, Location, etc.), you can add a CSV upload without changing the partition structure — the `Domain` filter continues to work.

5. What to Read Next

Ready to...	Read this
Implement the chosen architecture (step-by-step)	Multi-Agency Setup Guide (Setup layer)
Upload org data (CSV preparation and upload)	Data Onboarding Guide (Setup layer)
Set up the base Copilot analytics environment	Admin Setup Guide (Setup layer)
Understand the full data strategy (Entra ID vs. CSV)	Data Strategy Guide (Strategy layer)

Disclaimer

This guide is provided as a procedural reference. Always refer to Partitions in Viva Insights and Microsoft Copilot Dashboard on Microsoft Learn for the most up-to-date information.