

Microsoft 365 Copilot Analytics

Why Multi-Agency Analytics Is Hard

The architectural challenge of scoping analytics across agencies in a shared tenant

Document Property	Value
Section	Explain
Audience	CIO, Solution Architect
Version	4.0
Published	April 2026

When This Applies

Multi-agency analytics is a challenge when your organisation meets **all** of the following conditions:

- **Single M365 tenant** — all agencies share one tenant
- **Multiple independent agencies** (or business units, divisions, departments) that operate semi-autonomously
- **Each agency needs Copilot analytics** scoped to their own population

The fundamental tension: Microsoft 365 Copilot analytics is designed for one organisation per tenant. When multiple agencies share a tenant, there is no out-of-the-box way to guarantee that Agency A sees only its own data and never Agency B's. The Copilot Dashboard defaults to showing "Your company" — which means the entire tenant.

Two architectural approaches solve this, each with different trade-offs in data visibility, implementation effort, and ongoing maintenance.

Single-Agency Tenants

If your tenant serves only one agency (or your agencies don't need separate analytics scoping), none of this applies. Follow the standard Admin Setup Guide instead.

The Two Architectures

Default Priority is the simpler option. Each agency CEO naturally sees their team's data through the "Your group" scope on the Copilot Dashboard, which follows the Entra ID manager hierarchy. AutoCxO identification picks up the top leaders automatically. The trade-off: users *can* navigate to "Your company" and see tenant-wide data. If that cross-agency visibility is acceptable — and for many organisations it is — Default Priority is the right answer. Setup takes 2–4 hours.

Isolation is the harder option. It creates hard data boundaries via Viva Insights partitions — Agency A genuinely cannot access Agency B's data. This requires Viva feature access management (VFAM) lockdown via group-scoped policies (PowerShell), partition creation in the Viva Insights web app, analyst assignment per partition, and typically Power BI workspaces for report delivery. Setup takes 1–3 days. It exists for organisations with regulatory, contractual, or policy requirements that mandate data separation.

Within the Isolation architecture, there are two data paths depending on whether agencies share an email domain:

- **Path A** — shared domain (e.g., all users are `@government.gov`). Requires a CSV org data upload with a custom `Agency` attribute to distinguish agencies. Monthly re-upload needed.
- **Path B** — distinct domains (e.g., `@agencyA.gov` , `@agencyB.gov`). Uses the system-generated `Domain` attribute from Entra ID. No CSV needed for partitioning, though you may still want one for additional filter attributes.

Comparison at a Glance

Dimension	Default Priority	Isolation
Core principle	Agency CEOs see their team by default via manager hierarchy	Hard data boundaries via
Cross-agency visibility	Possible — users can navigate to "Your company"	Blocked — partition bour
Partitions needed	No	Yes — one per agency
Report delivery	Native Copilot Dashboard	Power BI workspaces (re
PowerShell required	No	Yes (VFAM lockdown onl
Setup effort	2–4 hours	1–3 days
Best for	Organisations where cross-agency visibility is acceptable	Organisations with manc

Ready to Decide?

If you understand the trade-offs above, the decision often comes down to one question: *must agencies be prevented from seeing each other's data?* If no, choose Default Priority. If yes, choose Isolation, then determine your data path based on whether agencies share an email domain.

[Use the Architecture Decision tool →](#)

Disclaimer

This guide is provided as a procedural reference. Always refer to Partitions in Viva Insights and Microsoft Copilot Dashboard on Microsoft Learn for the most up-to-date information.